

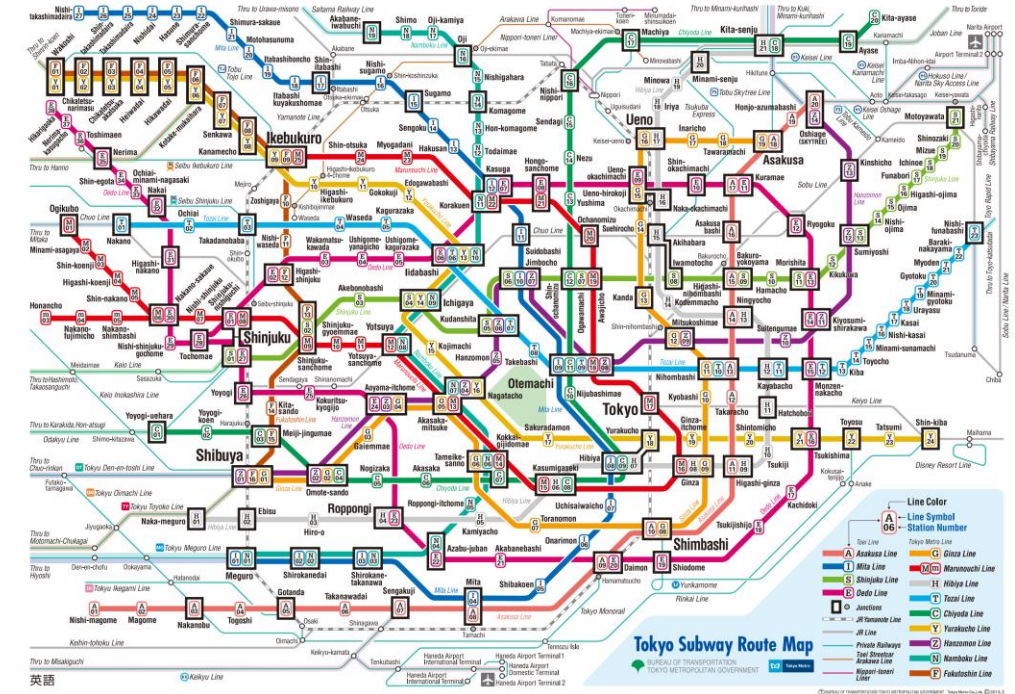


Programming in C/C++

- Graphs (2) – Shortest Paths -

Shortest Path

- Often, we don't only want to just find a path between two nodes, but we want to find the **shortest** path between two nodes v_1, v_2 in a graph
- Shortest path - path with smallest path weight
 - for unweighted graphs: number of edges in the path
 - for weighted graphs: sum of edge weights in the path





Dijkstra's Algorithm



Dijkstra's Algorithm


Basic idea:

- **Extend the current shortest path** starting a source and ending at node k by all its neighbors $n(k)$.
- The path weight for the neighboring node is the cost to k plus the weight for the edge to the neighbor $e(k, n(k))$
- For each neighbor $n(k)$, test whether the new link is shorter than the previous one
- Once destination node is reached output the path
- Required data structures:
 - Remember the best predecessor to trace back the best solution from end to source
 - Priority queue to find the node with currently shortest path to source node




Dijkstra's Algorithm

- Weighted Graph






Dijkstra's Algorithm



- Begin with S
- Update: shortest path to neighbors c and a have distance $< \infty$ -> store new distances and remember the predecessor ←




Dijkstra's Algorithm



- S already processed
- Add children a and c to priority queue to determine **next best** node

Queue: (c, 1) (a, 3)

Dijkstra's Algorithm




- Process c
- Extend path to all neighbor nodes

Queue: (a, 3)



Dijkstra's Algorithm




- Process c
- Extend path to all neighbor nodes – store new distances
- **update predecessors** ←
- update priority queue

Queue: (a, 2) (d, 5)

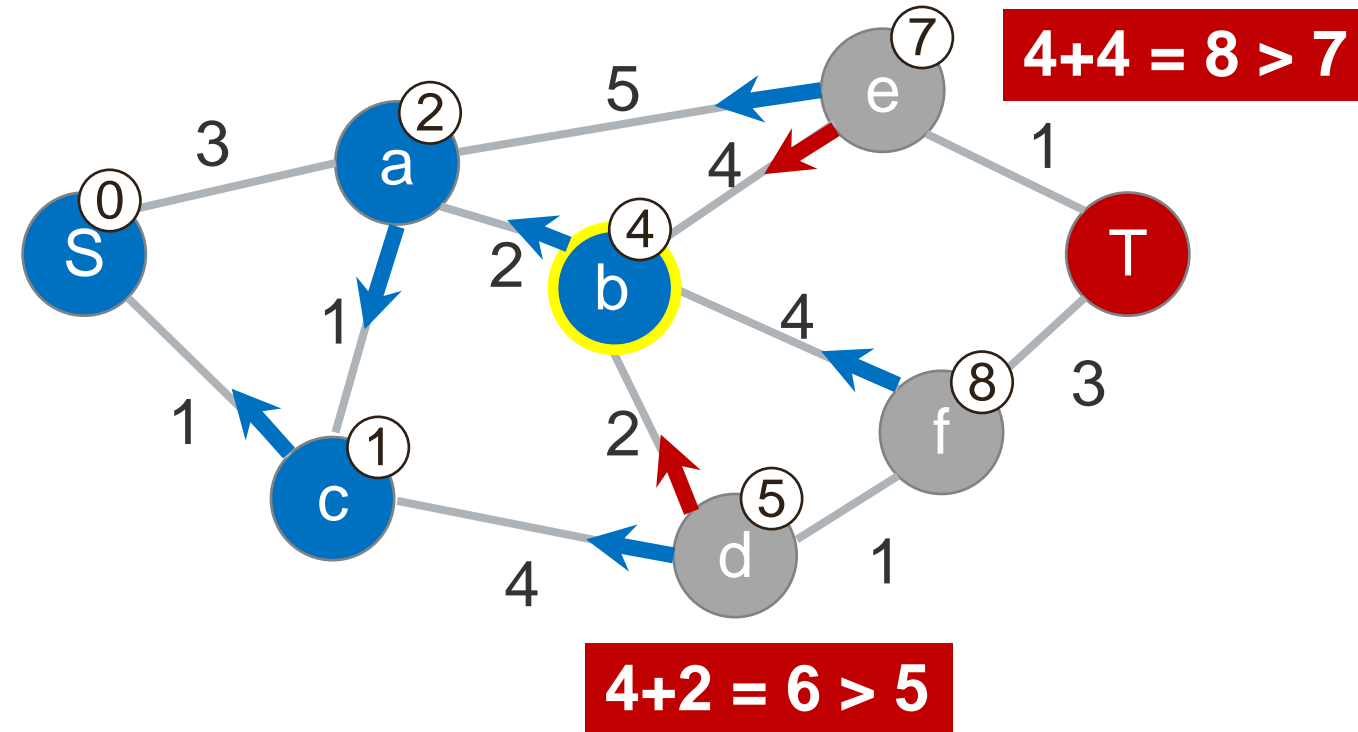


Dijkstra's Algorithm



- Process a
- Extend path to all neighbor nodes – store new distances
- **update predecessors**
- **update priority queue**

Queue : (b , 4) (d , 5) (e , 7)




- Process b
- Extend path to all neighbor nodes – store new distances
- (No better path to e or d, no update)

Queue: (d, 5) (e, 7) (f, 8)



Dijkstra's Algorithm




- Process d
- Extend path to all neighbor nodes – update if a cheaper path is found
- Update priority queue

Queue: (f, 6) (e, 7)



Dijkstra's Algorithm




- Process d
- Extend path to all neighbor nodes – update if a cheaper path is found
- Update priority queue

Queue: (e, 7) (T, 9)



Dijkstra's Algorithm




- Process e
- Extend path to all neighbor nodes – update if a cheaper path is found
- Update priority queue

Queue: (T, 9)



Dijkstra's Algorithm




- Process e
- Extend path to all neighbor nodes – update predecessor if a cheaper path is found
- Update priority queue

Queue: (T, 8)



Dijkstra's Algorithm



- Process T: Reached goal!
- Shortest path: unroll from T
- Queue is empty

Queue :


```

1  function Dijkstra(Graph, source):
2      for each vertex v in Graph:                // Initializations
3          dist[v] := infinity ;                  // Unknown distance function from
4                                                  // source to v
5          previous[v] := undefined ;             // Previous node in optimal path
6      end for                                     // from source

7
8      dist[source] := 0 ;                         // Distance from source to source
9      Q := the set of all nodes in Graph ;        // All nodes in the graph are
10                                                  // unoptimized - thus are in Q
11  while Q is not empty:                          // The main loop
12      u := vertex in Q with smallest distance in dist[] ;
13          // Source node in first case
14      remove u from Q ;
15      if dist[u] = infinity:
16          break ;                                // all remaining vertices are
17      end if                                     // inaccessible from source
18
19      for each neighbor v of u:                   // where v has not yet been
20                                                  // removed from Q.
21          alt := dist[u] + dist_between(u, v) ;
22          if alt < dist[v]:                       // Relax (u,v,a)
23              dist[v] := alt ;
24              previous[v] := u ;
25              decrease-key v in Q;               // Reorder v in the Queue
26          end if
27      end for
28  end while
29  return dist;
30 endfunction

```






Observation

The currently determined shortest subpath never becomes shorter by later processing of additional nodes

- Already visited nodes do not need to be looked at again

Priority queue - different implementations possible:

- Insertion of all nodes vs. insertion of just reachable nodes
- Heap vs. set (elements must also be changed)





Dijkstra using the BGL

- Boost provides free peer-reviewed portable C++ source libraries
- Many of the libraries later became part of the STL
- The BGL is a (not so modern anymore) library for Graph algorithms
- Most of the code is used to configure vertices and edges

```
// define vertex property to be a name and a color
using VertexPropertyType = property<vertex_name_t, std::string, property<vertex_color_t, default_color_type>>>;

// define edge property to contain an int weight and a color
using EdgePropertyType = property<edge_weight_t, int, property<edge_color_t, default_color_type>>>;

// graph will be an adjacency list with directed edges and our vertex and edge properties
using DirectedGraphType = adjacency_list<vecS, vecS, directedS, VertexPropertyType, EdgePropertyType>;

// vertex manipulation happens through handler
using VertexDescriptor = graph_traits<DirectedGraphType>::vertex_descriptor;

DirectedGraphType g;
VertexDescriptor a = boost::add_vertex(VertexPropertyType("a", white_color), g);
VertexDescriptor b = boost::add_vertex(VertexPropertyType("b", white_color), g);
VertexDescriptor c = boost::add_vertex(VertexPropertyType("c", white_color), g);
VertexDescriptor d = boost::add_vertex(VertexPropertyType("d", white_color), g);
VertexDescriptor e = boost::add_vertex(VertexPropertyType("e", white_color), g);
boost::add_edge(a, c, EdgePropertyType(1, black_color), g); add_edge(b, d, EdgePropertyType(1, black_color), g);
boost::add_edge(b, e, EdgePropertyType(2, black_color), g); add_edge(c, b, EdgePropertyType(5, black_color), g);
boost::add_edge(c, d, EdgePropertyType(10, black_color), g); add_edge(d, e, EdgePropertyType(4, black_color), g);
boost::add_edge(e, a, EdgePropertyType(3, black_color), g); add_edge(e, b, EdgePropertyType(7, black_color), g);
```



Dijkstra using the BGL

```
std::vector<int> distances(boost::num_vertices(g)); // Output for distances to each node
std::vector<VertexDescriptor> predMap(boost::num_vertices(g)); // Output for predecessors of each node in the shortest path tree result

// create predecessor/distance map
auto distanceMap = boost::predecessor_map(
    boost::make_iterator_property_map(predMap.begin(), boost::get(boost::vertex_index, g)).distance_map(
        boost::make_iterator_property_map(distances.begin(), boost::get(boost::vertex_index, g)));

// set start node
VertexDescriptor sourceV = a;

// run dijkstra
boost::dijkstra_shortest_paths(g, sourceV, distanceMap);

// set end node
VertexDescriptor destinationV = e;

int totalCost = distances[destinationV]; // read total path length

// trace path from end vertex to source vertex
std::vector<VertexDescriptor> path;
VertexDescriptor current = destinationV;
while (sourceV != current)
{
    path.push_back(current);
    current = predMap[current];
}
path.push_back(sourceV); // add source as last element to path

for (auto rit = path.rbegin(); rit != path.rend(); ++rit) // print path from source to destination (iterate using rbegin())
    std::cout << *rit << " -> ";
```

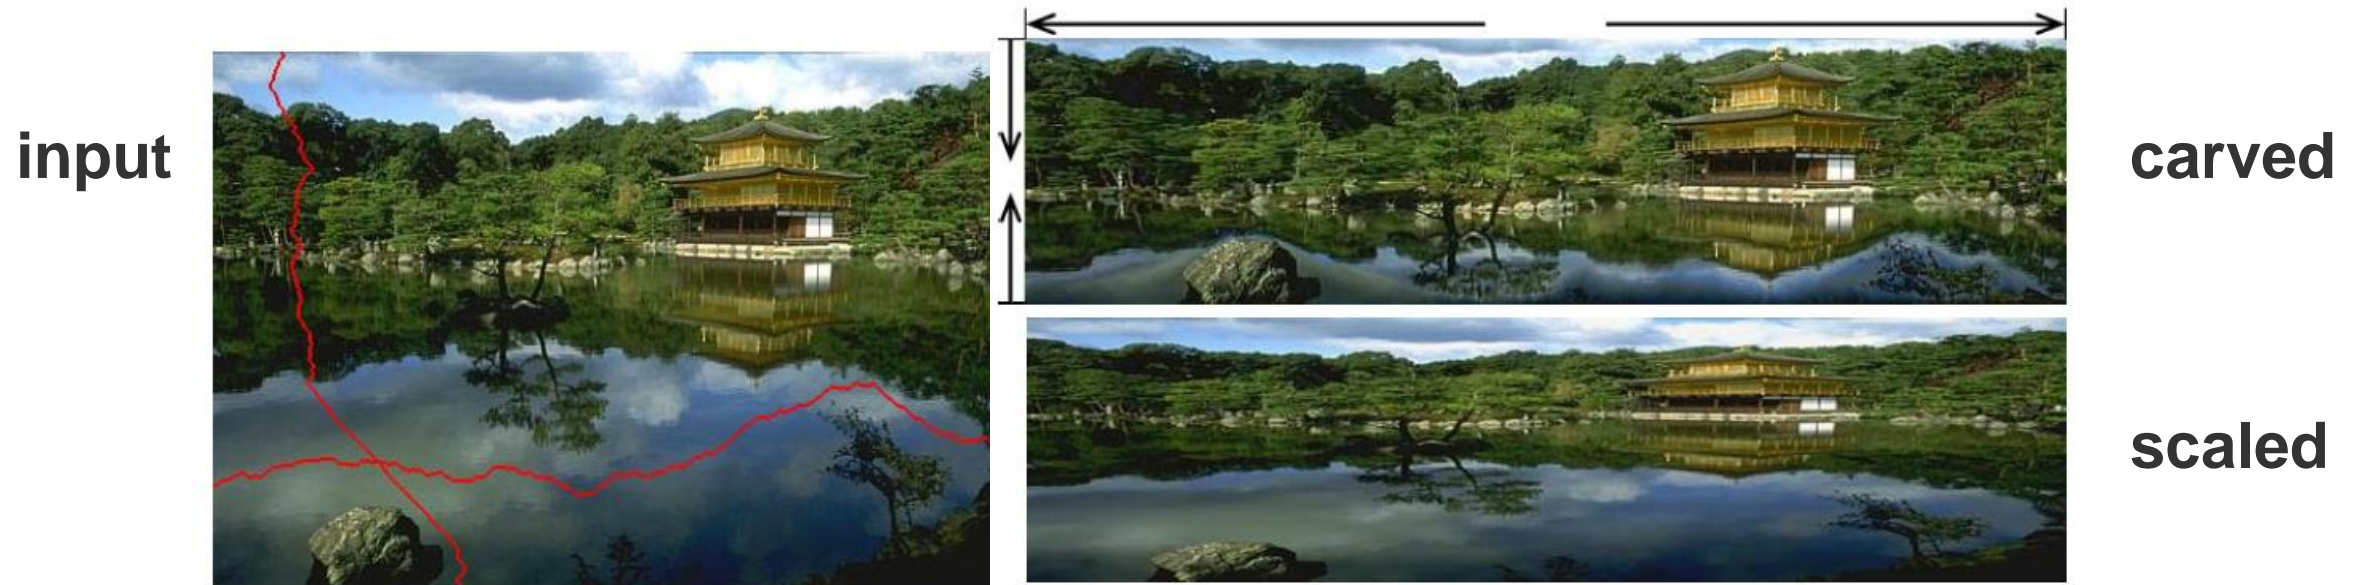


Dynamic Programming

Shortest paths on images

Seam Carving for Content-Aware Image Resizing

- [Shai Avidan, Ariel Shamir, ACM SIGGRAPH 2007]



- Goal: Change the size or aspect ratio of the image without distorting the content
- Repeatedly find the shortest path through the image

Seam Carving for Content-Aware Image Resizing



Seam Carving

- Find a path, from top to bottom (left/right), that destroys as little image content as possible
- Give importance of image content determined by changes to neighboring pixels

$$e = \left| \frac{\partial I}{\partial x} \right| + \left| \frac{\partial I}{\partial y} \right|$$

$$\frac{\partial I}{\partial x} \approx \frac{I(x+1, y) - I(x-1, y)}{2}$$


$$\frac{\partial I}{\partial y} \approx \frac{I(x, y+1) - I(x, y-1)}{2}$$

central differences

- source pixel (in top row) and destination pixel (in bottom row) are unknown
- Dijkstra for all point pairs would be too inefficient




Example: Shortest Path through an Image






Example: Shortest Path through an Image



- Graph implicitly given on image: e.g. virtual edges to three predecessors




Example: Shortest Path through an Image



- Graph implicitly given on image: e.g. virtual edges to three predecessors




Example: Shortest Path through an Image



- Graph implicitly given on image: e.g. virtual edges to three predecessors


Example: Shortest Path through an Image



- Graph implicitly given on image: e.g. virtual edges to three predecessors



Example: Shortest Path through an Image



- Graph implicitly given on image: e.g. virtual edges to three predecessors




Shortest Path with Dynamic Programming

- Path weight: sum of the node weights along the path

In each line

- For each point, calculate the shortest path to the row above it
 - From the **three** predecessors choose the one with the smallest path weight
 - Add node weight
- source(s) in the top row
- Remember predecessors
- In the last row, determine pixel with smallest path weight


Example: Shortest Path through an Image



- Trivial for first row: Shortest path of the pixel \rightarrow itself




Example: Shortest Path through an Image



- Second row:
 - Find the most favorable path from the three predecessors
 - Add node weight


Example: Shortest Path through an Image



- Second row:
 - Find the most favorable path from the three predecessors
 - Add node weight




Example: Shortest Path through an Image



- Second row:
 - Find the most favorable path from the three predecessors
 - Add node weight




Example: Shortest Path through an Image



- Second row:
 - Find the most favorable path from the three predecessors
 - Add node weight


Example: Shortest Path through an Image



- Second row:
 - Find the most favorable path from the three predecessors
 - Add node weight




Example: Shortest Path through an Image



- Second row:
 - Find the most favorable path from the three predecessors
 - Add node weight




Example: Shortest Path through an Image



- Second row:
 - Find the most favorable path from the three predecessors
 - Add node weight


Example: Shortest Path through an Image



- Third row:
 - Find the most favorable path from the three predecessors
 - Add node weight




Example: Shortest Path through an Image



- Third row:
 - Find the most favorable path from the three predecessors
 - Add node weight




Example: Shortest Path through an Image



- Third row:
 - Find the most favorable path from the three predecessors
 - Add node weight




Example: Shortest Path through an Image



- Fourth row:
 - Find the most favorable path from the three predecessors
 - Add node weight




Example: Shortest Path through an Image



- Fourth row:
 - Find the most favorable path from the three predecessors
 - Add node weight




Example: Shortest Path through an Image



- Fourth row:
 - Find the most favorable path from the three predecessors
 - Add node weight




Example: Shortest Path through an Image



- Last row:
 - Smallest accumulated path weight →
 - source of the shortest path




Example: Shortest Path through an Image




- Last row:
 - Smallest accumulated path weight →
 - source of the shortest path


Example: Seam Carving




- Remove all pixels along the shortest path



- Remove all pixels along the shortest path



- Remove all pixels along the shortest path



- Do ... repeat



Summary

- For large structured data, explicitly building a Graph is too much overhead.
- For structured data the graph is not explicitly built but e.g., implicitly given by image structure



A*-Algorithm



Recap Dijkstra's Algorithm – Shortest Path

The currently determined shortest subpath never becomes shorter by later processing of additional nodes

- Already visited nodes do not need to be looked at again

Priority queue:

- Insertion of all nodes vs. insertion of just reachable nodes
- Heap vs. set (elements must also be changed)

Properties of Dijkstra's algorithm:

- Similar to breadth-first search
- No targeted search
- Guarantees the shortest path
- Growth isotropically in all directions ("Sphere" from source)



A*-Algorithm

Basic idea:

- Mix depth and breadth-first search (if necessary)
- Like Dijkstra: expand the currently most favorable path
- Prioritization is based on an estimate of the **total cost** of a path **to the destination**:

$$w(p) = w(s, p) + w(p, t)$$

- Cost from source point to p known / propagated. $w(s, p)$
- Remaining cost from p to destination estimated with **heuristics** $w(p, t) \approx h(p)$



Condition on the Heuristic

The heuristic is monotonic if:

1. Actual costs are **never overestimated** by the heuristic
2. For each node p and successor p' must hold:

$$h(p) \leq w(p, p') + h(p')$$


(**triangle inequality**): estimated cost of a node p is smaller than the actual cost of the edge + estimated cost of that node

- If condition 2 is violated, the heuristic might still be valid
- **Example:** Euclidean distance (line of sight) in 2D maps is a monotonic heuristic for the actual travel distance between cities



A*-Algorithm – Heuristic


- Weighted graph





A*-Algorithm – Heuristic


- Weighted graph





A*-Algorithm – Heuristic


- Weighted graph





A*-Algorithm – Heuristic


- Weighted graph





A*-Algorithm – Heuristic


- Weighted graph





A*-Algorithm – Heuristic


- Weighted graph





A*-Algorithm – Heuristic


- Weighted graph








A*-Algorithm – Heuristic


- Weighted graph







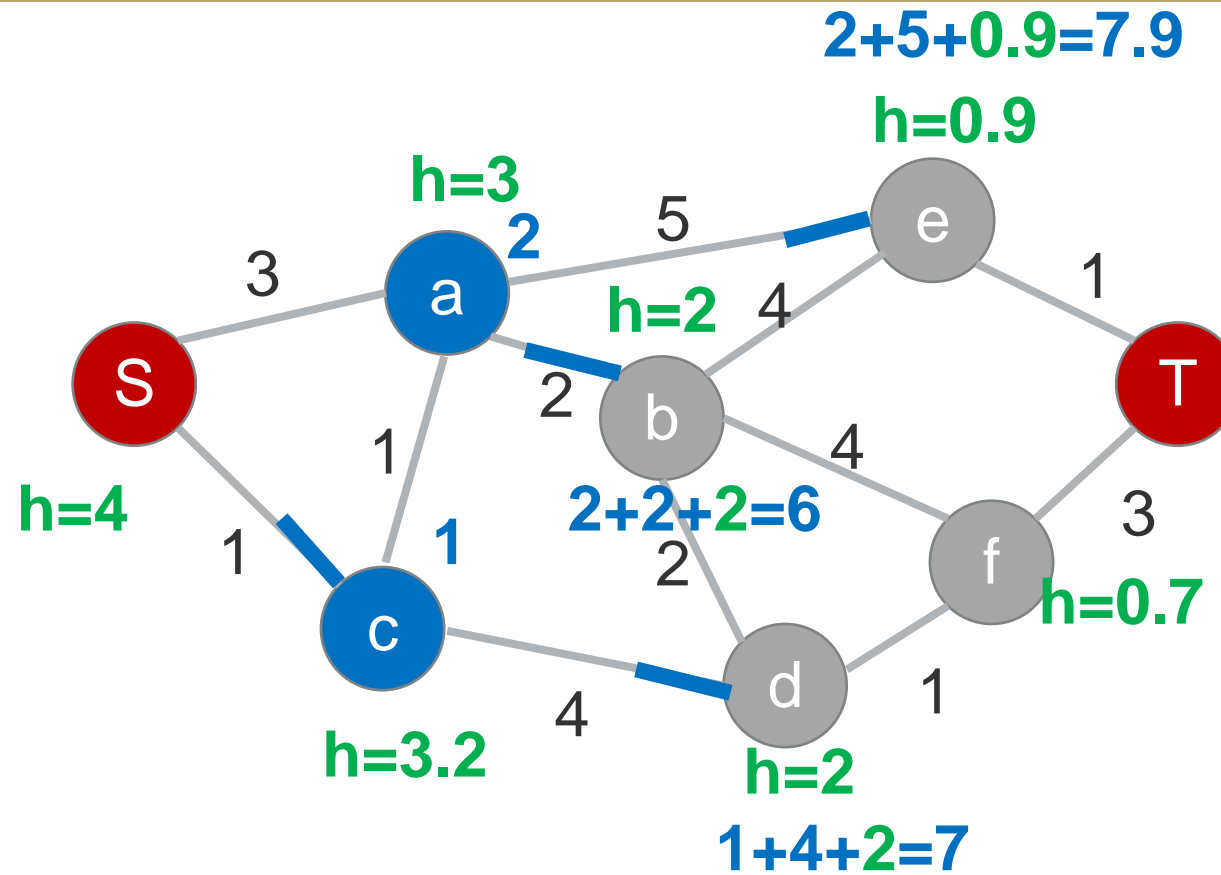


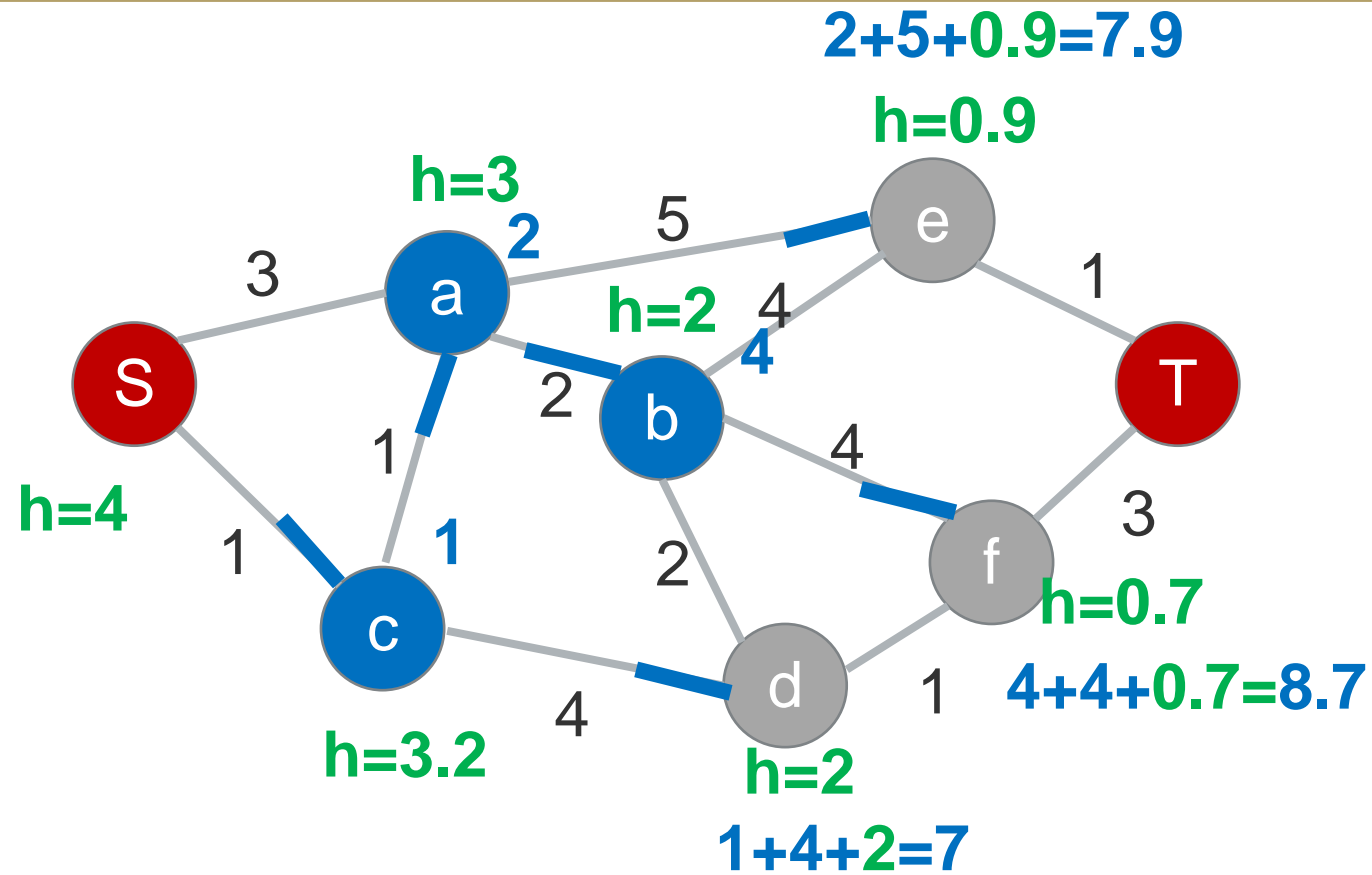








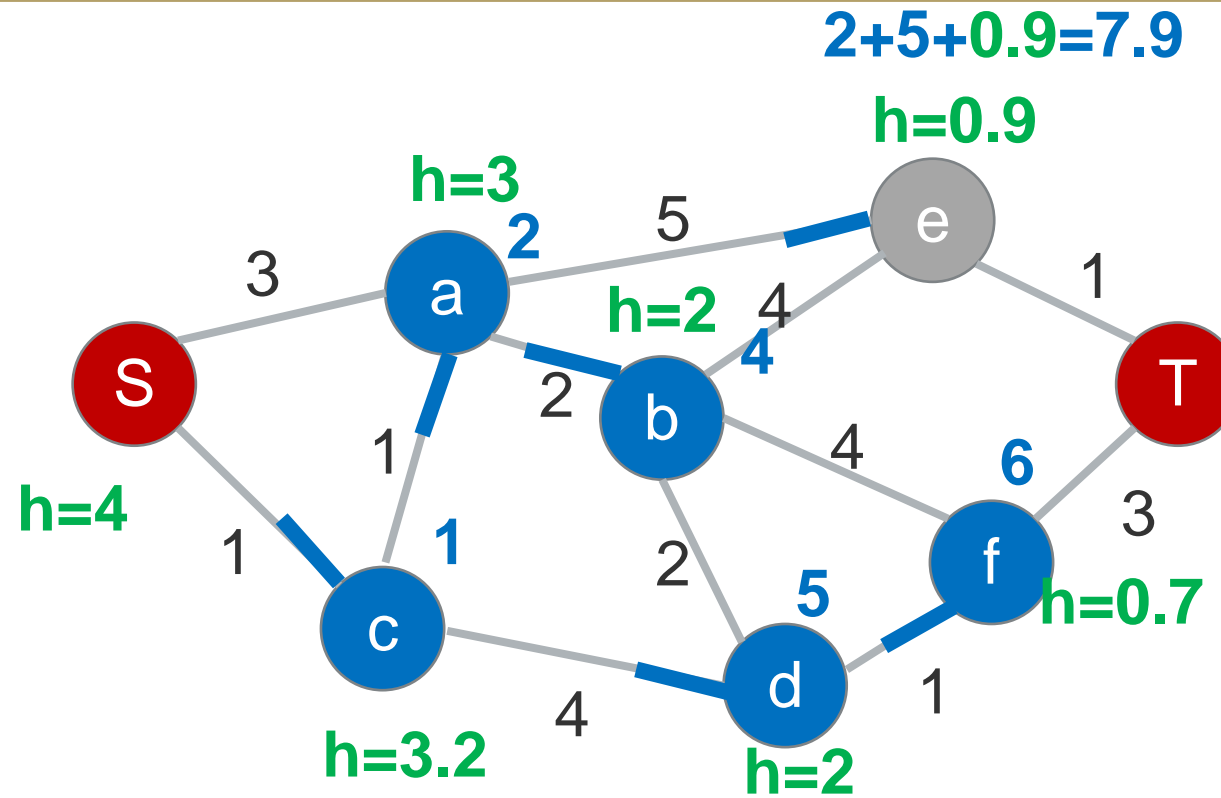


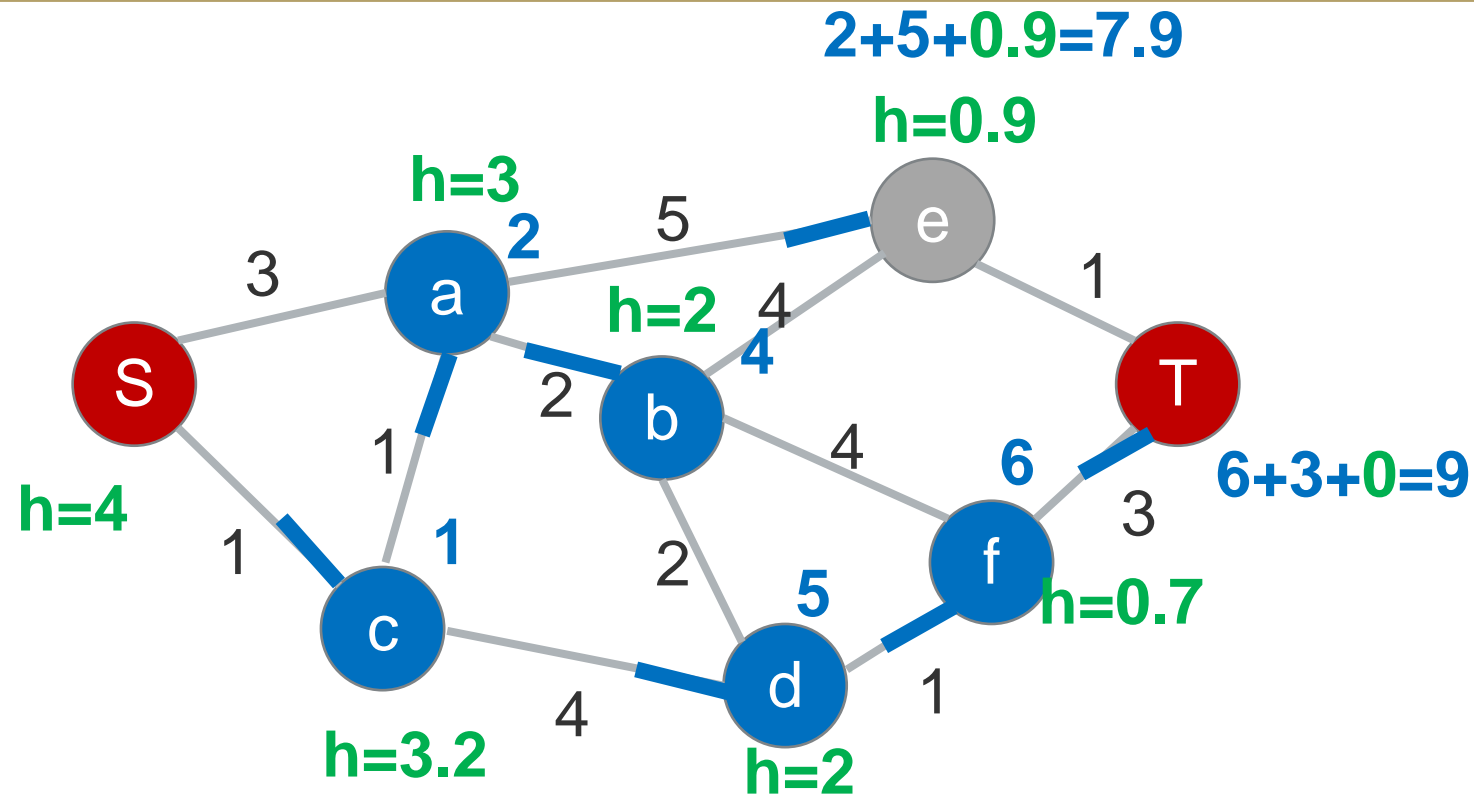





















A*-Algorithm - Discussion

The A* algorithm is

- **complete**: it finds a solution if a path exists
 - **optimal**: it finds a shortest path, if there are several shortest paths of the same length, it finds one of them
 - **optimally efficient**: for the given heuristic there is no other algorithm that finds the solution faster
-
- https://www.boost.org/doc/libs/1_80_0/libs/graph/example/astar-cities.cpp



A* visits significantly fewer sites





Summary

- Shortest path algorithm
 - Dijkstra
 - Dynamic programming
- A*-Algorithm
 - Shortest path with heuristics



Concluding the Lecture



Programming in C/C++

We covered a lot of ground in this intensive block course.

What didn't we cover?

A lot! ... and there is more coming in new C++ standards (theory) or becomes useable with better compiler support (practice).

We hope that the course prepared you to dive deeper into C++.

That you, if needed, are now able to read into the technical C++ documentation on topics we couldn't cover and apply what you learned in your projects.



Thanks

to the TAs for their help

to you for listening